

CMSC 510 – L09

Regularization Methods for Machine Learning



Instructor:
Dr. Tom Arodz

Recap: MLE for $P(y|x)$

- Let's say we have a probability distribution $P(y|x)$ of a certain shape, and the distribution is parameterized by a vector w , so $P(y|x)=P(y|x,w)$
- How to estimate w ?

- We have a training set S with m samples
- We could do maximum likelihood estimation of w

$$\max P(y_1, y_2, \dots, y_m | x_1, x_2, \dots, x_m, w)$$

- For which \mathbf{w} are the observed y 's for x 's most likely?
- We don't need to know anything about probability of x 's
 - we're not estimating $P(x,y)$, just $P(y|x)$

- Samples are i.i.d: they're independent, so:

$$P(y_1, y_2, \dots, y_m | x_1, x_2, \dots, x_m, w) = \prod_{i=1}^m P_w(y_i | x_i, w)$$

- ML estimate of \mathbf{w} :




\mathbf{w} that maximizes:

$$\max \prod_{i=1}^m P(y_i | x_i, w)$$

Recap: MLE / Logistic regression

- Under the assumption that class conditional probabilities depend on an unknown \mathbf{w} in this way:

$$P(y_i | x_i, w) = a(y_i w^T x_i) = \frac{1}{1 + e^{-y_i w^T x_i}}$$



$$P(+1 | x_i, w) \quad P(-1 | x_i, w) = 1 - P(+1 | x_i, w)$$

- Maximum likelihood estimate of w is:
- Solve this instead:

$$\min -\frac{1}{m} \ln \prod_{i=1}^m a(y_i w^T x_i)$$



- Or this:

$$\min \frac{1}{m} \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

Back to our view of LR


- Under the assumption that class conditional probabilities depend on an unknown \mathbf{w} in this way:

$$P(y_i | x_i, w) = \frac{1}{1 + e^{-y_i w^T x_i}}$$



$$P(+1 | x_i, w) \quad P(-1 | x_i, w) = 1 - P(+1 | x_i, w)$$

- Logistic regression minimizing the logistic loss: $\min \frac{1}{m} \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$



- is a *maximum likelihood* estimate of \mathbf{w} given the training set:

$$\max P(y_1, y_2, \dots, y_m | x_1, x_2, \dots, x_m, w)$$



Dealing with large weights

- Logistic regression is a *maximum likelihood* estimate of \mathbf{w} given the training set:

$$\arg \max_w P(y_1, y_2, \dots, y_m | x_1, x_2, \dots, x_m, w)$$

$$\arg \max_w \prod_{i=1}^m P(y_i | x_i, w)$$

- *Maximum a posteriori (MAP)* estimate of w is:

$$\arg \max_w P(y_1, y_2, \dots, y_m | x_1, x_2, \dots, x_m, w) P(w)$$

$$\arg \max_w P(w) \prod_{i=1}^m P(y_i | x_i, w)$$

- We take into account our estimate of probability $P(w)$ of each possible vector of weights w
 - Here, we can insert a belief that large w are unlikely: $P(\text{large } w)$ is small, $P(\text{small } w)$ is large

Dealing with large weights

- *Maximum a posteriori (MAP)* estimate of w is:

$$\arg \max_w P(y_1, y_2, \dots, y_m | x_1, x_2, \dots, x_m, w) P(w)$$
$$\arg \max_w P(w) \prod_{i=1}^m P(y_i | x_i, w)$$

- Applying the standard $\log()$ trick:

$$\arg \max_w \ln P(w) + \sum_{i=1}^m \ln P(y_i | x_i, w)$$

$$\arg \min_w -\ln P(w) + \sum_{i=1}^m -\ln P(y_i | x_i, w)$$

$$\arg \min_w \ln \frac{1}{P(w)} + \sum_{i=1}^m \ln \frac{1}{P(y_i | x_i, w)}$$

Dealing with large weights

- Logistic regression was a *maximum likelihood* estimate of \mathbf{w} given the training set, for a specific form of $P(y|x, \mathbf{w})$: $P(y_i | x_i, \mathbf{w}) = \frac{1}{1 + e^{-y_i \mathbf{w}^T x_i}}$

$$\arg \min_{\mathbf{w}} \sum_{i=1}^m \ln(1 + e^{-y_i \mathbf{w}^T x_i})$$

- *Maximum a posteriori (MAP)* estimate of \mathbf{w} for the same form of $P(y|x, \mathbf{w})$ is:

$$\arg \min_{\mathbf{w}} \left[\ln \frac{1}{P(\mathbf{w})} \right] + \sum_{i=1}^m \ln(1 + e^{-y_i \mathbf{w}^T x_i})$$

Dealing with large weights

- Let's assume: $P(y_i | x_i, w) = \frac{1}{1 + e^{-y_i w^T x_i}}$
- *Maximum a posteriori (MAP)* estimate of \mathbf{w} given the training set is:

$$\arg \max_w P(w) \prod_{i=1}^m P(y_i | x_i, w)$$


$$\arg \min_w \left[\ln \frac{1}{P(w)} \right] + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

- $P(w)$ = plausibility of this particular w
 - unless we have some evidence that large w are needed, we believe large w are unlikely to be truly necessary
 - More likely, they're artificial effect of correlated features
- $P(w)$ should be: small for large w ,
large for small w

Dealing with large weights

- $P(w)$ should be: small for large w , large for small w
- How to measure if w is large or small?

$$\begin{pmatrix} w_0 \\ w_1 \\ w_2 \end{pmatrix}$$



Dealing with large weights

- $P(w)$ should be: small for large w , large for small w
- How to measure if w is large or small?


L_p norm:

$$\|w\|_p = \left(\sum_{f=1}^F |w_f|^p \right)^{1/p}$$

large w = large value of $\|w\|$

we raise $\|w\|$ to power p for convenience

$$\|w\|_p^p = \left(\sum_{f=1}^F |w_f|^p \right)$$



Dealing with large weights

- Let's assume: $P(y_i | x_i, w) = \frac{1}{1 + e^{-y_i w^T x_i}}$
- *Maximum a posteriori (MAP)* estimate of **w**:

$$\arg \min_w \left[\ln \frac{1}{P(w)} \right] + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

- $P(w)$ should be: small for large L_p norm of w ,
large for small L_p norm of w

$$\|w\|_p^p = \left(\sum_{f=1}^F |w_f|^p \right)$$


- What form should $P(w)$ take?
 - $P(w) = f(\|w\|_p^p)$ where $0 \leq \|w\|_p^p \leq \infty$
 - What would be a convenient choice for $f()$?

Dealing with large weights

- Let's assume: $P(y_i | x_i, w) = \frac{1}{1 + e^{-y_i w^T x_i}}$
- *Maximum a posteriori (MAP)* estimate of **w**:

$$\arg \min_w \left[\ln \frac{1}{P(w)} \right] + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

- $P(w)$ should be: small for large w ,
large for small w



$$P(w) = \frac{1}{C} e^{-\frac{1}{C} \|w\|_p^p}$$

L_p norm


large w = large norm

$$\|w\|_p^p = \left(\sum_{f=1}^F |w_f|^p \right)$$


Dealing with large weights

- $P(w)$ should be: small for large w , large for small w

$$\|w\|_p^p = \left(\sum_{f=1}^F |w_f|^p \right)$$



$$P(w) = \frac{1}{C} e^{-\frac{1}{c} \|w\|_2^2}$$




$$P(w) = \frac{1}{C} e^{-\frac{1}{c} \|w\|_1}$$

Dealing with large weights

- Let's assume: $P(y_i | x_i, w) = \frac{1}{1 + e^{-y_i w^T x_i}}$
- *Maximum a posteriori (MAP)* estimate of **w**:

$$\arg \min_w \left[\ln \frac{1}{P(w)} \right] + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$


- Gaussian prior



Constant C is
standard deviation
or its equivalent

$$P(w) = \frac{1}{C} e^{-\frac{1}{C} \|w\|_2^2}$$

- Laplace prior





$$P(w) = \frac{1}{C} e^{-\frac{1}{C} \|w\|_1}$$

- Lower C \Rightarrow $P(w)$ concentrated more around 0
- Lower C \Rightarrow promotes smaller feature weights w



L_1 vs L_2 regularization

$$P(w) = \frac{1}{C} e^{-\frac{1}{C} \|w\|_p^p}$$

- $P(w)$ for L_2 regularization



- $P(w)$ for L_1 regularization





L_p regularization

- L_p regularization

$$\arg \min_w \left[\ln \frac{1}{P(w)} \right] + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

- What does the first term expand to?

- We have:

$$P(w) = \frac{1}{C} e^{-\frac{1}{C} \|w\|_p^p}$$

- So:

$$\ln \frac{1}{P(w)} = \frac{1}{C} \|w\|_p^p + \ln C$$

Regularization

$$\ln \frac{1}{P(w)} = \frac{1}{C} \|w\|_p^p + \ln C$$


- Logistic regression with L_p regularization

$$\arg \min_w \left[\ln \frac{1}{P(w)} \right] + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

$\rightarrow \arg \min_w \frac{1}{C} \|w\|_p^p + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$


- We're adding a penalty term $\frac{1}{C} \|w\|_p^p$

- The penalty term looks like this:



$$\|w\|_2^2$$

or this:



$$\|w\|_1$$

$$\|w\|_p^p = \left(\sum_{f=1}^F |w_f|^p \right)$$

Regularization


- Logistic regression with L_p regularization

$$\arg \min_w \frac{1}{C} \|w\|_p^p + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

- We're adding a penalty term $\frac{1}{C} \|w\|_p^p$
- Algorithm minimizes:
empirical risk of $h()$ + penalty for complexity of $h()$


Minimum:

$w \rightarrow 56.3318,$
 $w_0 \rightarrow -48.3115$



$$\sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$


+



$$\|w\|_2^2$$

No penalty over w_0

=

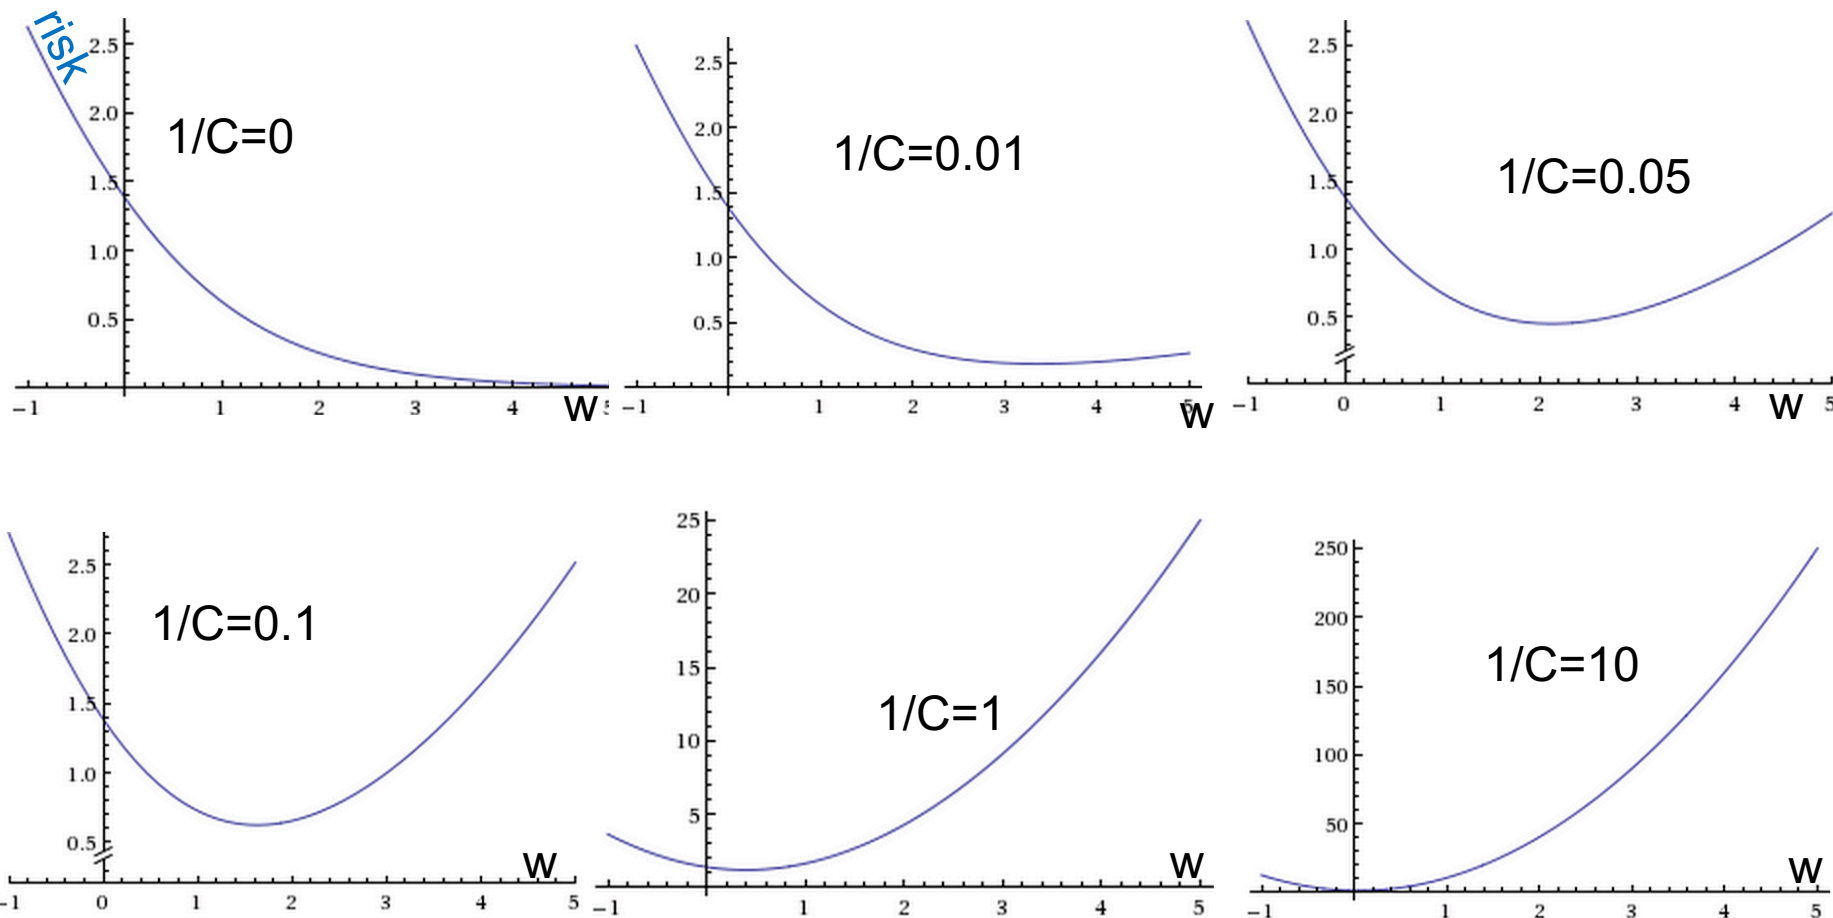


Minimum:
 $w \rightarrow 0.647365,$
 $w_0 \rightarrow -1.14344$

L₂ Regularization

$$\arg \min_w \frac{1}{C} \|w\|_p^p + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

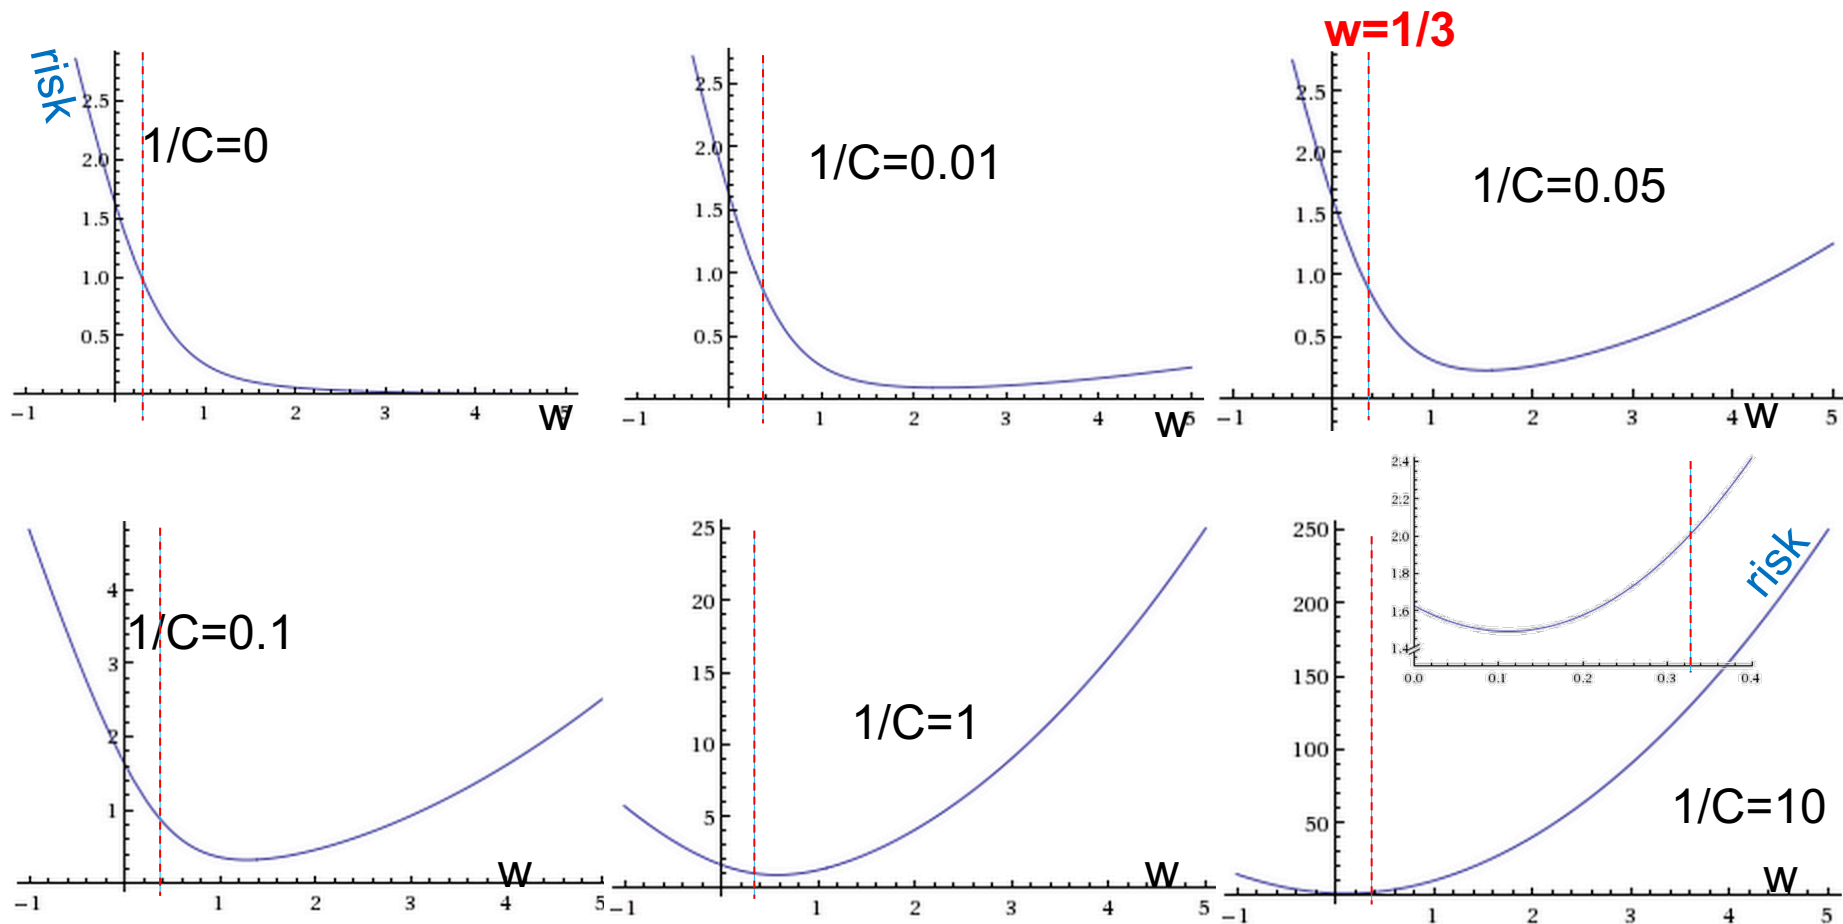
- Logistic regression with L₂ regularization
- Higher 1/C => higher penalty for large norm of w
 - Example: two points (x=1,y=1), (x=-1,y=-1), w₀=0 (fixed value)



L₂ Regularization

$$\arg \min_w \frac{1}{C} \|w\|_p^p + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

- Example: two points $(x=3, y=1)$, $(x=-1, y=-1)$, $w_0=-1$ (fixed value)
 - $w x + w_0 = 3 * w + (-1) > 0 \Rightarrow \mathbf{w > 1/3}$ (condition for correct prediction)
 - $w x + w_0 = -1 * w + (-1) < 0 \Rightarrow w > -1$ (condition for correct prediction)
- Too much regularization is bad: $1/C=10 \Rightarrow w=0.11 \Rightarrow$ wrong prediction



L₂ Regularization

- Logistic regression:


$$\arg \min_w \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

- Solved through gradient descent:

$$w_{t+1} = w_t - \frac{\partial \left(\sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i}) \right)}{\partial w}$$

$$w_{t+1} = w_t - \sum_{i=1}^m \frac{\partial \ln(1 + e^{-y_i w^T x_i})}{\partial w}$$

- Weights may grow and grow,
gradient never going to 0



Minimum:
w -> 56.3318,
w₀ -> -48.3115

L₂ Regularization

- Logistic regression with L₂ regularization:

$$\arg \min_w \frac{1}{C} \|w\|_2^2 + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i})$$

- Solved through gradient descent:

$$w_{t+1} = w_t - \frac{\partial \left(\frac{1}{C} \|w\|_2^2 + \sum_{i=1}^m \ln(1 + e^{-y_i w^T x_i}) \right)}{\partial w}$$
$$w_{t+1} = w_t - \frac{2}{C} w_t - \sum_{i=1}^m \frac{\partial \ln(1 + e^{-y_i w^T x_i})}{\partial w}$$

- **Weight decay!**

- In each iteration, old weights are reduced by a fraction
 - Before we add something new, from gradient
- Weights don't grow to be large



Coding classification methods

Thankfully, we do not have to act like in HW1

In the last couple of years, a number of libraries for automating gradient descent became popular

Tensorflow, PyTorch


Automatic differentiation

Tensorflow/PyTorch are libraries for performing calculations and derivatives on a *computational graph*

Example: $y = f(x_1, x_2) = \ln(x_1) + x_1 x_2 - \sin(x_2)$

We want to compute y and dy/dx_1 for:


$$(x_1, x_2) = (2, 5)$$



Automatic differentiation

In machine learning, we rarely want to calculate gradient of loss w.r.t. feature values x_i

Instead, we will have model weights (w, b) as the starting variables for the computational graph (training samples x, y would just be constants somewhere in the graph)




Automatic differentiation (AD)

Forward Primal Trace

$v_{-1} = x_1$	$= 2$
$v_0 = x_2$	$= 5$
<hr/>	
$v_1 = \ln v_{-1}$	$= \ln 2$
$v_2 = v_{-1} \times v_0$	$= 2 \times 5$
$v_3 = \sin v_0$	$= \sin 5$
$v_4 = v_1 + v_2$	$= 0.693 + 10$
$v_5 = v_4 - v_3$	$= 10.693 + 0.959$
<hr/>	
$y = v_5$	$= 11.652$

$$y = f(x_1, x_2) = \ln(x_1) + x_1 x_2 - \sin(x_2) \quad (x_1, x_2) = (2, 5)$$



AD

Dot over a variable represents a derivative
Over what? Below, over x_1
There will be a similar trace for x_2


Forward Primal Trace

$v_{-1} = x_1$	$= 2$
$v_0 = x_2$	$= 5$
<hr/>	
$v_1 = \ln v_{-1}$	$= \ln 2$
$v_2 = v_{-1} \times v_0$	$= 2 \times 5$
$v_3 = \sin v_0$	$= \sin 5$
$v_4 = v_1 + v_2$	$= 0.693 + 10$
$v_5 = v_4 - v_3$	$= 10.693 + 0.959$
<hr/>	
$y = v_5$	$= 11.652$

Forward Tangent (Derivative) Trace

$\dot{v}_{-1} = \dot{x}_1$	$= 1$
$\dot{v}_0 = \dot{x}_2$	$= 0$
<hr/>	
$\dot{v}_1 = \dot{v}_{-1}/v_{-1}$	$= 1/2$
$\dot{v}_2 = \dot{v}_{-1} \times v_0 + \dot{v}_0 \times v_{-1}$	$= 1 \times 5 + 0 \times 2$
$\dot{v}_3 = \dot{v}_0 \times \cos v_0$	$= 0 \times \cos 5$
$\dot{v}_4 = \dot{v}_1 + \dot{v}_2$	$= 0.5 + 5$
$\dot{v}_5 = \dot{v}_4 - \dot{v}_3$	$= 5.5 - 0$
<hr/>	
$\dot{y} = \dot{v}_5$	$= 5.5$

$$y = f(x_1, x_2) = \ln(x_1) + x_1 x_2 - \sin(x_2) \quad (x_1, x_2) = (2, 5)$$





AD

Dot over a variable represents a derivative
Over what? Below, over x_1
There will be a similar trace for x_2

Forward Primal Trace

$v_{-1} = x_1$	$= 2$
$v_0 = x_2$	$= 5$
<hr/>	
$v_1 = \ln v_{-1}$	$= \ln 2$
$v_2 = v_{-1} \times v_0$	$= 2 \times 5$
$v_3 = \sin v_0$	$= \sin 5$
$v_4 = v_1 + v_2$	$= 0.693 + 10$
$v_5 = v_4 - v_3$	$= 10.693 + 0.959$
<hr/>	
$y = v_5$	$= 11.652$

Forward Tangent (Derivative) Trace

$\dot{v}_{-1} = \dot{x}_1$	$= 1$
$\dot{v}_0 = \dot{x}_2$	$= 0$
<hr/>	
$\dot{v}_1 = \dot{v}_{-1}/v_{-1}$	$= 1/2$
$\dot{v}_2 = \dot{v}_{-1} \times v_0 + \dot{v}_0 \times v_{-1}$	$= 1 \times 5 + 0 \times 2$
$\dot{v}_3 = \dot{v}_0 \times \cos v_0$	$= 0 \times \cos 5$
$\dot{v}_4 = \dot{v}_1 + \dot{v}_2$	$= 0.5 + 5$
$\dot{v}_5 = \dot{v}_4 - \dot{v}_3$	$= 5.5 - 0$
<hr/>	
$\dot{y} = \dot{v}_5$	$= 5.5$

$$y = f(x_1, x_2) = \ln(x_1) + x_1 x_2 - \sin(x_2) \quad (x_1, x_2) = (2, 5)$$

We do not get derivatives in a form of mathematical formulas

We just get the value of the derivative for specific x_1, x_2

That's what we need for gradient descent!

Forward Primal Trace

$v_{-1} = x_1$	$= 2$
$v_0 = x_2$	$= 5$
<hr/>	
$v_1 = \ln v_{-1}$	$= \ln 2$
$v_2 = v_{-1} \times v_0$	$= 2 \times 5$
$v_3 = \sin v_0$	$= \sin 5$
$v_4 = v_1 + v_2$	$= 0.693 + 10$
$v_5 = v_4 - v_3$	$= 10.693 + 0.959$
<hr/>	
$y = v_5$	$= 11.652$


Forward Tangent (Derivative) Trace

$\dot{v}_{-1} = \dot{x}_1$	$= 1$
$\dot{v}_0 = \dot{x}_2$	$= 0$
<hr/>	
$\dot{v}_1 = \dot{v}_{-1} / v_{-1}$	$= 1/2$
$\dot{v}_2 = \dot{v}_{-1} \times v_0 + \dot{v}_0 \times v_{-1}$	$= 1 \times 5 + 0 \times 2$
$\dot{v}_3 = \dot{v}_0 \times \cos v_0$	$= 0 \times \cos 5$
$\dot{v}_4 = \dot{v}_1 + \dot{v}_2$	$= 0.5 + 5$
$\dot{v}_5 = \dot{v}_4 - \dot{v}_3$	$= 5.5 - 0$
<hr/>	
$\dot{y} = \dot{v}_5$	$= 5.5$

Above, we had forward differentiation

Calculate $d v_i / d x$ for increasing i

Until we get to $d v_5 / d x = d y / d x$




AD

There is an alternative way, closer to backpropagation

Calculate dy/dv_i for decreasing i

Until we get to $dy/dv_0 = dy/dx$



Forward Primal Trace

$v_{-1} = x_1$	$= 2$
$v_0 = x_2$	$= 5$
<hr/>	
$v_1 = \ln v_{-1}$	$= \ln 2$
$v_2 = v_{-1} \times v_0$	$= 2 \times 5$
<hr/>	
$v_3 = \sin v_0$	$= \sin 5$
$v_4 = v_1 + v_2$	$= 0.693 + 10$
<hr/>	
$v_5 = v_4 - v_3$	$= 10.693 + 0.959$
<hr/>	
$y = v_5$	$= 11.652$

Reverse Adjoint (Derivative) Trace

alternative way (faster)

$\bar{x}_1 = \bar{v}_{-1}$	$= 5.5$
$\bar{x}_2 = \bar{v}_0$	$= 1.716$
<hr/>	
$\bar{v}_{-1} = \bar{v}_{-1} + \bar{v}_1 \frac{\partial v_1}{\partial v_{-1}} = \bar{v}_{-1} + \bar{v}_1 / v_{-1}$	$= 5.5$
$\bar{v}_0 = \bar{v}_0 + \bar{v}_2 \frac{\partial v_2}{\partial v_0} = \bar{v}_0 + \bar{v}_2 \times v_{-1}$	$= 1.716$
$\bar{v}_{-1} = \bar{v}_2 \frac{\partial v_2}{\partial v_{-1}} = \bar{v}_2 \times v_0$	$= 5$
$\bar{v}_0 = \bar{v}_3 \frac{\partial v_3}{\partial v_0} = \bar{v}_3 \times \cos v_0$	$= -0.284$
$\bar{v}_2 = \bar{v}_4 \frac{\partial v_4}{\partial v_2} = \bar{v}_4 \times 1$	$= 1$
$\bar{v}_1 = \bar{v}_4 \frac{\partial v_4}{\partial v_1} = \bar{v}_4 \times 1$	$= 1$
$\bar{v}_3 = \bar{v}_5 \frac{\partial v_5}{\partial v_3} = \bar{v}_5 \times (-1)$	$= -1$
$\bar{v}_4 = \bar{v}_5 \frac{\partial v_5}{\partial v_4} = \bar{v}_5 \times 1$	$= 1$
<hr/>	
$\bar{v}_5 = \bar{y}$	$= 1$