

CMSC 510 – L20

Regularization Methods for Machine Learning




Part 20a: BatchNorm

Instructor:
Dr. Tom Arodz

Training deep nets


- Historically, deep networks (with many layers) were difficult to train




- A lot of effort in recent years went into training deep nets easier
 - ReLU instead of sigmoid
 - Unsupervised pre-training
 - Normalization techniques

Batch normalization

- A standard deep net has layers, where each layer is
 - a linear transformation Wx , where x is input from previous layer
 - a nonlinear activation function acting on Wx , e.g. $\text{ReLU}(Wx)$




- We can look at the training from the perspective of a **single layer** – as if the other layers were constant



Batch normalization


- We can look at the training from the perspective of a single layer – as if the other layers were constant



- Layer k will learn how to transform its input (output of layers $1, \dots, k-1$) into something that minimizes the "new loss" (layers $k+1, \dots, L$ + original loss).

Batch normalization


- We can look at the training from the perspective of a single layer – as if the other layers were constant



- If layers $1, \dots, k-1$ are not constant (are trained), the distribution of input to layer k changes all the time
- If layers $k+1, \dots, L$ are not constant (are trained), the “new loss” also changes all the time
- If we fix the first problem, inductively we also fix the second problem

Batch normalization


- We can look at the training from the perspective of a single layer – as if the other layers were constant



- If layers $1, \dots, k-1$ are not constant (are trained), the distribution of input to layer k changes all the time
- We can't fix it, really!
 - We do want to early layers to learn, i.e., change what they're producing
- But we can fix some general property of what it produces

Batch normalization

- We can look at the training from the perspective of a single layer – as if the other layers were constant



- We can fix some general property of what it produces
- Normalization:
 - Make some statistic (e.g. mean) of the outputs of a layer constant, even if the actual output vectors change during training



Batch normalization

- Normalization:

- Make some statistic (e.g. mean, or std.dev) of the outputs of a layer constant, even if the actual output vectors change during training

- Output of a layer with 6 neurons, on a batch of 4 samples

- a 4x6 matrix WX

$$\begin{bmatrix} 1 & -3 & 4 & -2 & 5 & 4 \\ 2 & -6 & 9 & -1 & 8 & 2 \\ 2 & -6 & 9 & -1 & 9 & 7 \\ -1 & 3 & -4 & 2 & -5 & -4 \end{bmatrix}$$

- Possible options:

- Normalize rows (separately each sample, across the neurons)

- To add up to 1
- Or to have 0 mean

- Normalize columns (separately each neuron, across the batch):

- To add up to 1
- To have 0 mean



Batch normalization

- Normalization:

- Make some statistic (e.g. mean, or std.dev) of the outputs of a layer constant, even if the actual output vectors change during training

- BatchNorm

- Normalize columns (separately each neuron, across the batch):
 - Squares add up to 1, i.e. Std.Dev=1
 - Mean = 0

$$BN(y_j)^{(b)} = \gamma \cdot \left(\frac{y_j^{(b)} - \mu(y_j)}{\sigma(y_j)} \right) + \beta$$


- Add scale and offset parameters (so Mean is beta, Std.Dev. is gamma)




BatchNorm

$$BN(y_j)^{(b)} = \gamma \cdot \left(\frac{y_j^{(b)} - \mu(y_j)}{\sigma(y_j)} \right) + \beta$$

Standard Network




Adding a BatchNorm layer (between weights and activation function)




Batch normalization

- These are normal tensorflow / pytorch computations in the computational graph
 - i.e. chain rule applies to the operation of calculating mean/std.dev



(a) Vanilla Network



(b) Vanilla Network + BatchNorm Layer

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$


$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

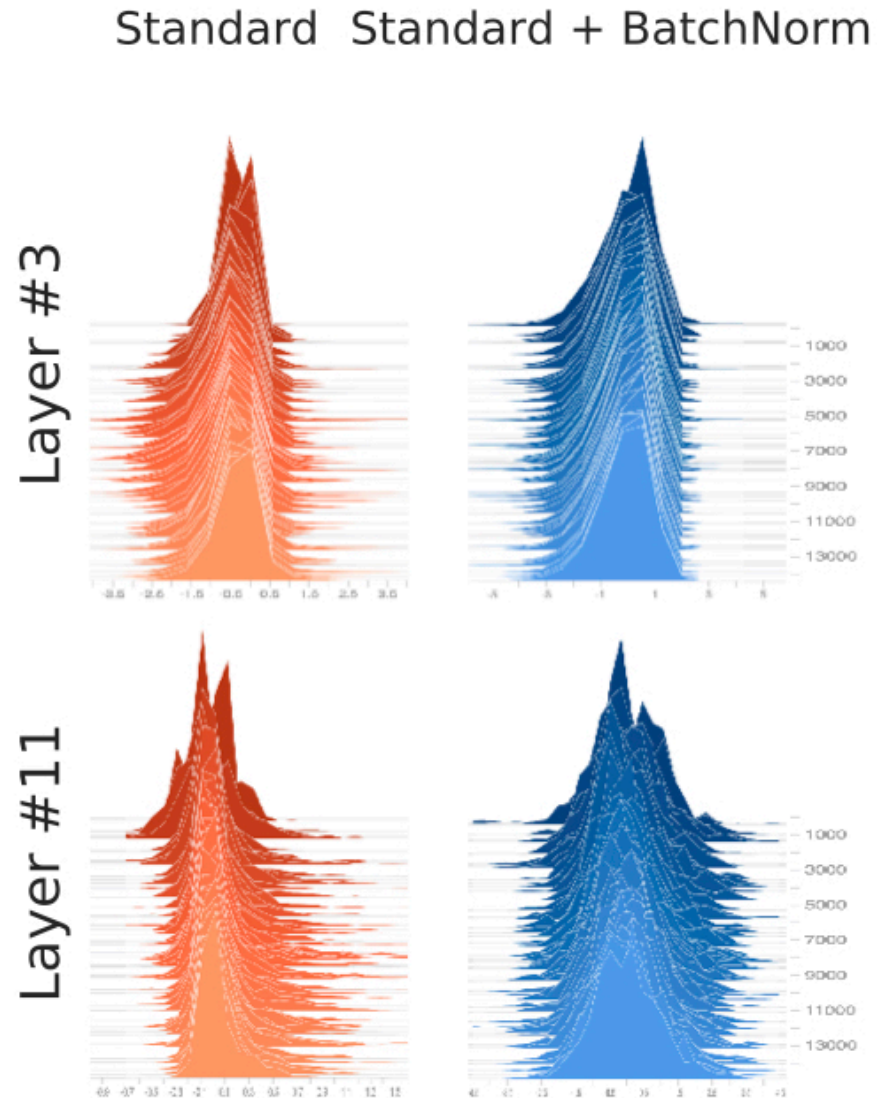
Batch normalization

- Does BatchNorm help?



Batch normalization

- How does BatchNorm help?
- BatchNorm is supposed to help with “internal covariate shift”
 - Normalize columns (separately each neuron, across the batch):
 - Squares add up to 1, i.e. $\text{Std.Dev}=1$
 - Mean = 0
- But even without BatchNorm we don't see much instability in the distributions





Batch normalization

- A three-layer neural network with ReLU activation is:
 - $Y = \text{ReLU}(W_3 \text{ReLU}(W_2 \text{ReLU}(W_1 X)))$
- Expanding $\text{ReLU}(wx) = \max(0, wx)$, we see either wx , or 0
 - $\text{ReLU}(w_2 \text{ReLU}(w_1 x))$ can be $w_2 * w_1 * x$
- Jointly over three layers, we see terms like $z = w_{3ij} * w_{2kl} * w_{1mn} * x$
- The derivative of z over w_{1mn} is $w_{3ij} * w_{2kl} * x$
- The derivative can quickly get large even if individual w 's are not that much larger than 1
- Or can get very small if individual w 's are close to zero




Batch normalization

- A neural network with ReLU activation is essentially:
 - $Y = \text{ReLU}(W_3 \text{ReLU}(W_2 \text{ReLU}(W_1 X)))$
- We see terms like $w_{3ij} * w_{2kl} * w_{1mn} * x$
- The derivative over w_{2kl} is $w_{3ij} * w_{1mn} * x$
- The derivative can quickly get large even if individual w 's are not that much larger than 1
- Or can get very small if individual w 's are close to zero
- Vanishing/exploding gradient!

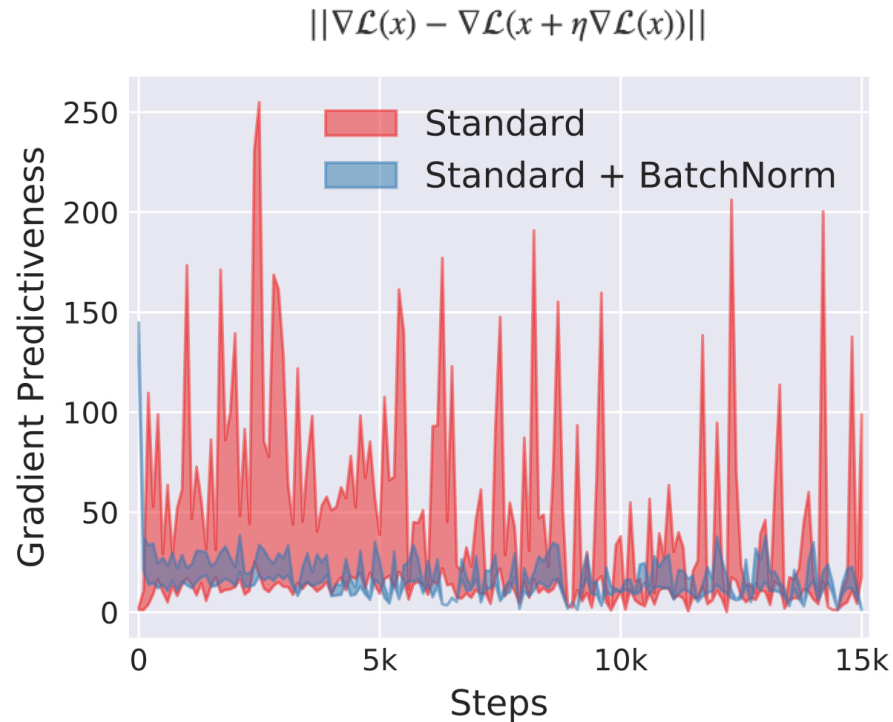
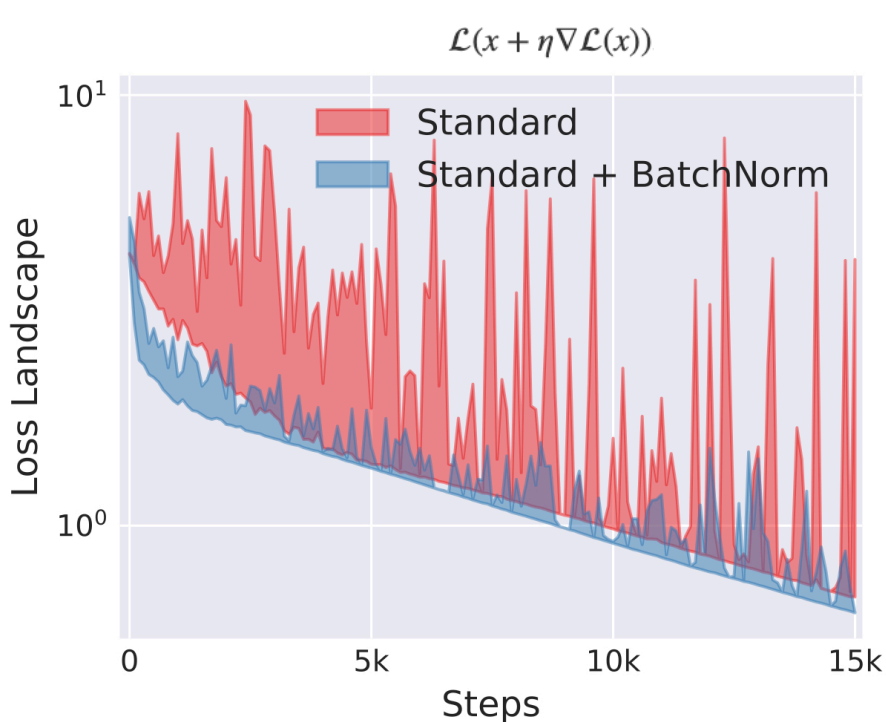
Batch normalization

- Vanishing gradients slow learning
- Exploding gradients prevent learning




Batch normalization

- The derivative over w_{2kl} is $w_{3ij} * w_{1mn} * x$
 $\sim w_{3ij} * \text{activation}(\text{prev. layer})$
- Solution: make activations “just right”
 - E.g. make them roughly follow a Gaussian with 0-mean, unit norm
- Help keep loss stable:



Normalizations

- Possible normalization options:
 - Layer norm: Normalize rows (separately each sample, across the neurons)
 - Batch norm: Normalize columns (separately each neuron, across the batch)
- In ConvNets we also have a third axis (channels / “colors”)
 - Like Layer norm, but within each channel separately



Multiple filters,
each with its own
set of weights

