

CMSC 510

Regularization Methods for Machine Learning




Graph Neural Networks

Instructor:
Dr. Tom Arodz




CNN vs GNN - similarities

Convolutional Neural Networks

- Input is provided on a regular lattice (pixels in a grid)



- At every graph node (pixel) we have a vector of node features
 - In the input layer: RGB colors
 - In subsequent layers: filter activations




Graph Neural Networks

- Input is provided on a general graph (typically not regular)



- At every graph node we have a vector of node features
 - E.g. it's a social network, a node is a person, features describe that person
 - Or a chemical molecule, each node has some properties of that atom




CNN vs GNN - differences

Convolutional Neural Networks

- Subsequent layers SOMEWHAT preserve the underlying lattice (pixel grid)

- But: often there is pooling to reduce the lattice size




- The output is unrelated to input graph – it's vector of class probabilities

Graph Neural Networks

- Subsequent layers TYPICALLY preserve the underlying graph

- Often there is no pooling, the structure is preserved all they way till the end



- Often, the output is the same graph as input – we're predicting some missing information (e.g. class) for each node



- Each node is like a "sample", some have class, some don't, like in semi-supervised learning

Graph Neural Networks

- Key intuition

- Information about a node come


- From its features
 - But also from features of its neighbors in the graph



Graph Neural Networks

Regularization

- Like in semi-supervised learning
 - Leverage information from neighbors in a graph
 - Leverage information from unlabeled samples
- Like in CNNs
 - "Weight sharing" – single set of weights applied at different nodes



Initial "layer 0" embeddings are equal to node features

$$\mathbf{h}_v^0 = \mathbf{x}_v$$

previous layer embedding of v

$$\mathbf{h}_v^k = \sigma \left(\mathbf{W}_k \sum_{u \in N(v)} \frac{\mathbf{h}_u^{k-1}}{|N(v)|} + \mathbf{B}_k \mathbf{h}_v^{k-1} \right), \quad \forall k > 0$$

kth layer embedding of v

non-linearity (e.g., ReLU or tanh)

average of neighbor's previous layer embeddings

Graph Neural Networks

Challenges

- In CNNs, filters can have a fixed structure (e.g. 3x3)
 - Graphs are irregular – how to account for that?
- Graphs (e.g. social networks) are often small-world
 - After a few hops, everything is connected to everything

